

DM 3 - convex optimisation

Titouan Le Breton

Question 1

We introduce the auxiliary variable

$$z = Xw - y \iff z + y - Xw = 0,$$

so the problem becomes

$$\min_{w,z} \frac{1}{2}\|z\|_2^2 + \lambda\|w\|_1 \quad \text{s.t.} \quad z + y - Xw = 0.$$

The Lagrangian :

$$L(w, z, \mu) = \frac{1}{2}\|z\|_2^2 + \lambda\|w\|_1 + \mu^\top(z + y - Xw) = \mu^\top y + \left(\frac{1}{2}\|z\|_2^2 + \mu^\top z\right) + (\lambda\|w\|_1 - \mu^\top Xw).$$

Thus the dual function is

$$g(\mu) = \inf_{w,z} L(w, z, \mu) = \mu^\top y + \inf_z \left(\frac{1}{2}\|z\|_2^2 + \mu^\top z\right) + \inf_w (\lambda\|w\|_1 - \mu^\top Xw).$$

For the z -term, note that

$$\frac{1}{2}\|z\|_2^2 + \mu^\top z = \sum_{i=1}^n \left(\frac{1}{2}z_i^2 + \mu_i z_i\right)$$

so the infimum is attained at $z = -\mu$ and equals $-\frac{1}{2}\|\mu\|_2^2$.

For the w -term, write $a_j = \mu^\top x_j$ where x_j is the j -th column of X , so

$$\lambda\|w\|_1 - \mu^\top Xw = \lambda \sum_{j=1}^d |w_j| - \sum_{j=1}^d a_j w_j.$$

If $w_j = t$, with $t > 0$ and $w_i = 0$ if $i \neq j$ one has:

$$\lambda\|w\|_1 - \mu^\top Xw = \lambda t - a_j t = (\lambda - a_j)t.$$

If $\lambda - a_j < 0$ then $(\lambda - a_j)t \rightarrow -\infty$ as $t \rightarrow +\infty$, so the function is unbounded below in this direction.

If $w_j = -t$, with $t > 0$ and $w_i = 0$ if $i \neq j$, one has:

$$\lambda\|w\|_1 - \mu^\top Xw = \lambda t - a_j(-t) = (\lambda + a_j)t.$$

If $\lambda + a_j < 0$ then $(\lambda + a_j)t \rightarrow -\infty$ as $t \rightarrow +\infty$ and the function is unbounded below in this direction.

Therefore, for the objective to be bounded below for all such directions, we must have

$$\lambda - a_j \geq 0 \quad \text{and} \quad \lambda + a_j \geq 0,$$

i.e.

$$-\lambda \leq a_j \leq \lambda \quad \forall j,$$

which is exactly

$$\|X^\top \mu\|_\infty = \max_j |a_j| \leq \lambda.$$

Now assume that this condition holds, i.e. $\|X^\top \mu\|_\infty \leq \lambda$. Since $\|X^\top \mu\|_\infty \leq \lambda$, we have $|a_j| \leq \lambda$ for all j . Write

$$\lambda|w_j| - a_j w_j = \lambda\left(|w_j| - \frac{a_j}{\lambda} w_j\right).$$

Since $|\frac{a_j}{\lambda}| \leq 1$, we deduce that for all w , the function is positive, with it attaining a minimum of 0 for $w = 0$.

Hence the dual problem is

$$\max_{\mu \in \mathbb{R}^n} \mu^\top y - \frac{1}{2} \|\mu\|_2^2 \quad \text{s.t.} \quad \|X^\top \mu\|_\infty \leq \lambda.$$

Rewriting as a minimization yields

$$\min_{\mu \in \mathbb{R}^n} \frac{1}{2} \|\mu\|_2^2 - y^\top \mu \quad \text{s.t.} \quad X^\top \mu \leq \lambda \mathbf{1}_d, \quad -X^\top \mu \leq \lambda \mathbf{1}_d.$$

This fits the quadratic program $\min v^\top Qv + p^\top v$ with

$$Q = \frac{1}{2} I_n, \quad p = -y,$$

and constraint matrix

$$A = \begin{pmatrix} X^\top \\ -X^\top \end{pmatrix}, \quad b = \lambda \begin{pmatrix} \mathbf{1}_d \\ \mathbf{1}_d \end{pmatrix}.$$

$$\min_{v \in \mathbb{R}^n} v^\top \left(\frac{1}{2} I_n \right) v - y^\top v \quad \text{s.t.} \quad \begin{pmatrix} X^\top \\ -X^\top \end{pmatrix} v \leq \lambda \begin{pmatrix} \mathbf{1}_d \\ \mathbf{1}_d \end{pmatrix}.$$

Question 2

Given the quadratic objective

$$f_0(x) = x^\top Qx + p^\top x,$$

and the inequality constraints $Ax \leq b$, the barrier-augmented function for a parameter $\tau > 0$ is

$$F_\tau(x) = \tau(x^\top Qx + p^\top x) - \sum_{i=1}^m \log(b_i - a_i^\top x),$$

where a_i^\top is the i -th row of A .

The gradient splits as

$$\nabla F_\tau(x) = \tau \nabla(x^\top Qx + p^\top x) + \sum_{i=1}^m \nabla[-\log(b_i - a_i^\top x)].$$

For the quadratic part,

$$\nabla(x^\top Qx) = 2Qx, \quad \nabla(p^\top x) = p,$$

so

$$\tau \nabla(x^\top Qx + p^\top x) = \tau(2Qx + p).$$

For the barrier term,

$$\nabla \left[-\log(b_i - a_i^\top x) \right] = \frac{a_i}{b_i - a_i^\top x}.$$

Thus the full gradient is

$$\nabla F_\tau(x) = \tau(2Qx + p) + \sum_{i=1}^m \frac{a_i}{b_i - a_i^\top x}.$$

For the Hessian, differentiate again:

$$\nabla \left(\frac{a_i}{b_i - a_i^\top x} \right) = \frac{a_i a_i^\top}{(b_i - a_i^\top x)^2}.$$

Hence the full Hessian is

$$\nabla^2 F_\tau(x) = 2\tau Q + \sum_{i=1}^m \frac{a_i a_i^\top}{(b_i - a_i^\top x)^2}.$$

Question 3

In Figure 1 and Figure 2, the effect of different values of μ is illustrated. Figure 1 corresponds to 10 datapoints in dimension 100, while Figure 2 corresponds to 100 datapoints in dimension 1000. In both cases, the results are similar.

The left side of the figures show how convergence changes for different values of μ . We see that higher μ values lead to significantly faster convergence, which is no surprise. The L2 distance matrix to the right side of the plot compares the final estimation of the w values. They are very similar (considering that these distance values are for high dimension), notably showing that increasing μ does not hinder performance.

As such, a high value of μ like 50 or 100 to get both low compute time and good results is the best configuration.




Figure 1: Convergence of barr_method and distance matrix of the final w for different values of μ for $n = 10$ and $d = 100$




Figure 2: Convergence of barr_method and distance matrix of the final w for different values of μ for $n = 100$ and $d = 100$

Importing packages

```
In [1]: import numpy as np  
import matplotlib.pyplot as plt
```

I. Centering step

Functions for gradient and hessian computation

```
In [2]: def compute_gradient_ft(t, x, Q, p, A, b):  
    # t the scalar that characterises the function we are trying to minimise  
    # x of dimension (n), the point in which we are evaluating the function  
    # Q of dimension (n,n)  
    # p of dimension (n)  
    # A of dimension (m,n)  
    # b of dimension (m)  
  
    m = b.size  
    n = x.size  
    left_side = t*(2*Q @ x + p)  
    right_side = np.zeros(n)  
    for i in range(m):  
        #     print((b[i] - A[i] @ x).shape)  
        #     print(f"A shape: {A.shape}")  
        #     print(f"x shape: {x.shape}")  
        #     print(f"b shape: {b.shape}")  
        right_side += A[i] / (b[i] - A[i] @ x)  
  
    return left_side + right_side  
  
def compute_hessian_ft(t, x, Q, A, b):  
    # t the scalar that characterises the function we are trying to minimise  
    # x of dimension (n), the point in which we are evaluating the function  
    # Q of dimension (n,n)  
    # A of dimension (m,n)  
    # b of dimension (m)  
  
    m = b.size  
    n = x.size  
    left_side = 2 * t * Q  
    right_side = np.zeros((n,n))  
    for i in range(m):  
        right_side += (np.outer(A[i], A[i])) / (b[i] - A[i] @ x)**2  
  
    return left_side + right_side
```

Additonal utility functions

```
In [3]: # returns bool to check if the stop condition of the newton algorithm is met
def newton_dercrement(grad, delta_x, eps):
    return - grad @ delta_x / 2 <= eps

def compute_delta_x(grad, hessian):
    return np.linalg.solve(hessian, -grad)

def backtracking(ft, x, delta_x, grad, alpha, beta): # returns the optimal x con
    # x of shape (n)
    # delta_x of shape (n)
    # grad of shape (n)
    step = 1
    while ft(x + step * delta_x) >= ft(x) + alpha * step * grad @ delta_x:
        step = beta * step
    return x + step * delta_x
```

Centring step function

```
In [4]: def centering_step(Q, p, A, b, t, v0, eps):
    # initialisation
    v_running = v0
    list_v_points = []
    count = 0 # to ensure there is at least one iteration
    alpha = 0.3
    beta = 0.5

    while True:
        # do the matrix computations
        grad = compute_gradient_ft(t, v_running, Q, p, A, b)
        hessian = compute_hessian_ft(t, v_running, Q, A, b)
        delta_v = compute_delta_x(grad, hessian)

        # check for the stop condition
        if newton_dercrement(grad, delta_v, eps) and count > 0:
            break

        # define the f_t function so it can be used in backtracking
        def ft(x):
            s = b - A @ x
            if np.any(s <= 0): # value at +infinity if out of bounds
                return np.inf

            right_side = t * (x @ Q @ x + p @ x)
            left_side = -np.sum(np.log(s))
            return right_side + left_side

        # backtracking to upgrade v with the optimal step
        v_running = backtracking(ft, v_running, delta_v, grad, alpha, beta)
        list_v_points.append(v_running)
        count += 1

    return list_v_points

# dimensions
n = 10
d = 2

#LASSO hyperparameter
lambda_val = 10
```

```

# primal matrix sampling
X = np.random.randn(n, d)
y = np.random.randn(n)

# dual variables reformating
Q = np.eye(n) / 2
p = -y
A = np.concatenate([X.T, -X.T], axis=0)
b = lambda_val * np.ones(2 * d)

# interior points and Newton hyperparametres
v0 = np.zeros(n)
eps_newton = 0.01

t = 1
return_array = centering_step(Q, p, A, b, t, v0, eps_newton)

```

Log-barrier method (interior points)

```

In [5]: def barr_method(Q, p, A, b, v0, eps_barr=0.1, eps_newton=0.1, mu = 1.2):
    # set the parameters needed
    m = b.size
    t = 1

    # run Newton once
    list_v_barr = [v0]
    v_running = centering_step(Q, p, A, b, t, v0, eps_newton)[-1]
    list_v_barr.append(v_running)

    # iterate through the algorithm as long as m/t < eps is not verified
    while eps_barr <= m / t:
        t = mu * t
        v_running = centering_step(Q, p, A, b, t, v0, eps_newton)[-1]
        list_v_barr.append(v_running)

    return list_v_barr

# dimensions
n = 100
d = 30

#LASSO hyperparameter
lambda_val = 10

def generate_variables(n, d, lambda_val):
    # primal matrix sampling
    X = np.random.randn(n, d)
    y = np.random.randn(n)

    # dual variables reformating
    Q = np.eye(n) / 2
    p = -y
    A = np.concatenate([X.T, -X.T], axis=0)
    b = lambda_val * np.ones(2 * d)

    v0 = np.zeros(n)

```

```

    return Q, p, A, b, v0

Q, p, A, b, v0 = generate_variables(n, d, lambda_val)

# interior points and Newton hyperparametres

eps_barr = 0.1
mu = 1.2

list_v_barr = barr_method(Q, p, A, b, v0, eps_barr=eps_barr, mu=mu)

```

Running and plotting results

```

In [12]: def f(x, Q, p):
    return x @ Q @ x + p @ x

def plot_all_in_one(n, d, mus, lambda_val, eps_barr=0.1, eps_newton=0.1):
    fig, (ax_traj, ax_mat) = plt.subplots(1, 2, figsize=(12, 5))

    final_vectors = [] # store list_v_barr[-1] for each mu

    # Left: convergence curves
    for mu in mus:
        Q, p, A, b, v0 = generate_variables(n, d, lambda_val)
        list_v_barr = barr_method(Q, p, A, b, v0,
                                  eps_barr=eps_barr,
                                  eps_newton=eps_newton,
                                  mu=mu)

        x_final = list_v_barr[-1]
        final_vectors.append(x_final)

        f_opt = f(x_final, Q, p)
        f_diff = [f(v, Q, p) - f_opt for v in list_v_barr]
        ax_traj.plot(f_diff, label=f"mu={mu}")

    ax_traj.set_title(f"Convergence (n={n}, d={d})")
    ax_traj.set_xlabel("iterations")
    ax_traj.set_ylabel("distance to f_opt")
    ax_traj.legend()

    # right: MSE matrix between final vectors
    X = np.vstack(final_vectors) # shape (k, d)
    diff = X[:, None, :] - X[None, :, :]
    mse_mat = np.mean(diff**2, axis=2) # shape (k, k)

    im = ax_mat.imshow(mse_mat, interpolation='nearest')
    plt.colorbar(im, ax=ax_mat)
    ax_mat.set_title("L2 distance between w estimation for different mu")
    ax_mat.set_xticks(range(len(mus)))
    ax_mat.set_yticks(range(len(mus)))
    ax_mat.set_xticklabels(mus)
    ax_mat.set_yticklabels(mus)
    ax_mat.set_xlabel("mu")
    ax_mat.set_ylabel("mu")

```

```

plt.tight_layout()
plt.show()


```

In [13]:

```

lambda_val = 10
n = 10
d = int(1e2)
mu_list = [2, 15, 25, 50, 100]
plot_all_in_one(n, d, mu_list, lambda_val)

```




In [9]:

```

n = 100
d = int(1e3)
mu_list = [2, 15, 25, 50, 100]
plot_all_in_one(n, d, mu_list)

```



observation: increasing μ reduces the number of iterations and, considering the high-dimension w is in, the mse between (w_μ) is small. So a high value of μ like 50 or 100 should be selected so as to have fewer steps and faster computation.